

# DATENBANK-ARCHITEKTUR FÜR FORTGESCHRITTENE

## Performanceoptimierung

Dani Schnider

# Performance Tuning Tasks

## Performance Planning

- Scalability
- System Architecture
- Application Design
- Data Model
- Testing
- etc.

## Instance Tuning

- Memory Allocation
- I/O Balancing
- Database Configuration
- System Configuration
- etc.

## SQL Tuning

- Indexing
- Partitioning
- Rewrite SQL Statements
- Gathering Statistics
- Hints
- etc.

See [Oracle® Database Performance Tuning Guide, Chapter 1, Performance Tuning Overview](#)

# QUERY OPTIMIZER

**Performanceoptimierung**

# Query Optimizer

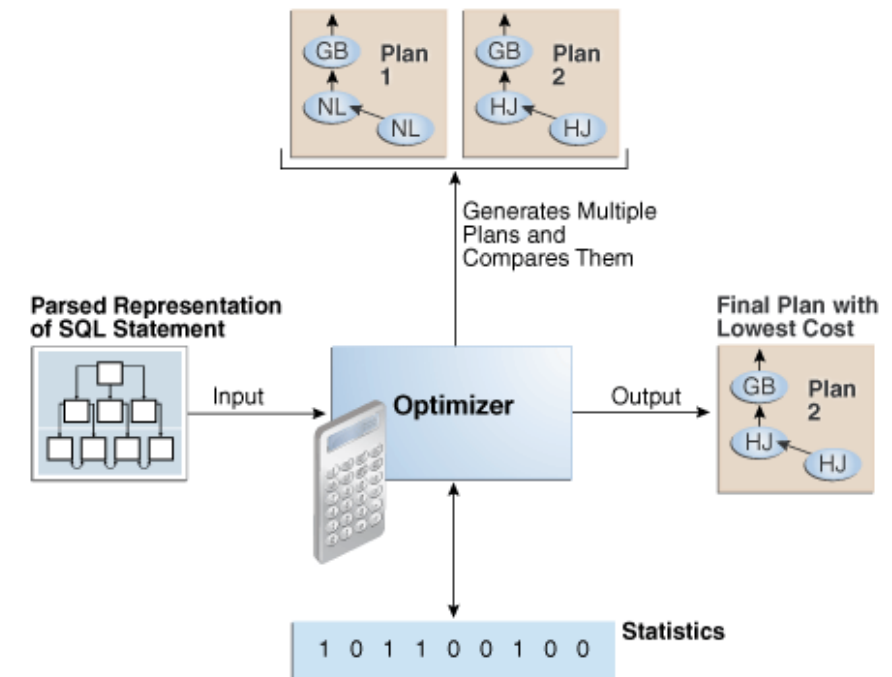
## Ziel:

- Finde effizientesten Weg zur Ausführung einer SQL-Befehls

## Vorgehensweise:

- Für jeden SQL-Befehl werden folgende Schritte durchgeführt:
  1. Generierung von mehreren möglichen Execution Plans
  2. Bewertung jedes Execution Plans mit Kosten\*
  3. Ausführung des Execution Plans mit den geringsten Kosten
- Als Basis für die Kostenberechnung werden Optimizer-Statistiken verwendet

\* Die Kosten sind ein geschätzter Wert für die I/O- und CPU-Ressourcen die zur Ausführung des Plans benötigt werden



# Execution Plan

Execution Plan legt fest, wie der SQL-Befehl ausgeführt wird:

- Join-Reihenfolge der Tabellen
- Join-Methode pro Tabelle
  - Nested Loops Join
  - Merge Join
  - Hash Join
- Zugriffs-Methode pro Tabelle
  - Full Table Scan
  - Index Scan

-----		
Id	Operation	Name
-----		
0	SELECT STATEMENT	
* 1	HASH JOIN	
* 2	TABLE ACCESS FULL	PRODUCTS
3	NESTED LOOPS	
4	NESTED LOOPS	
5	NESTED LOOPS	
6	TABLE ACCESS BY INDEX ROWID BATCHED	CUSTOMERS
* 7	INDEX RANGE SCAN	CUST_LAST_NAME
8	TABLE ACCESS BY INDEX ROWID BATCHED	ORDERS
* 9	INDEX RANGE SCAN	ORD_CUST_ID
* 10	INDEX RANGE SCAN	ORDI_ORDER_ID
11	TABLE ACCESS BY INDEX ROWID	ORDER_ITEMS
-----		

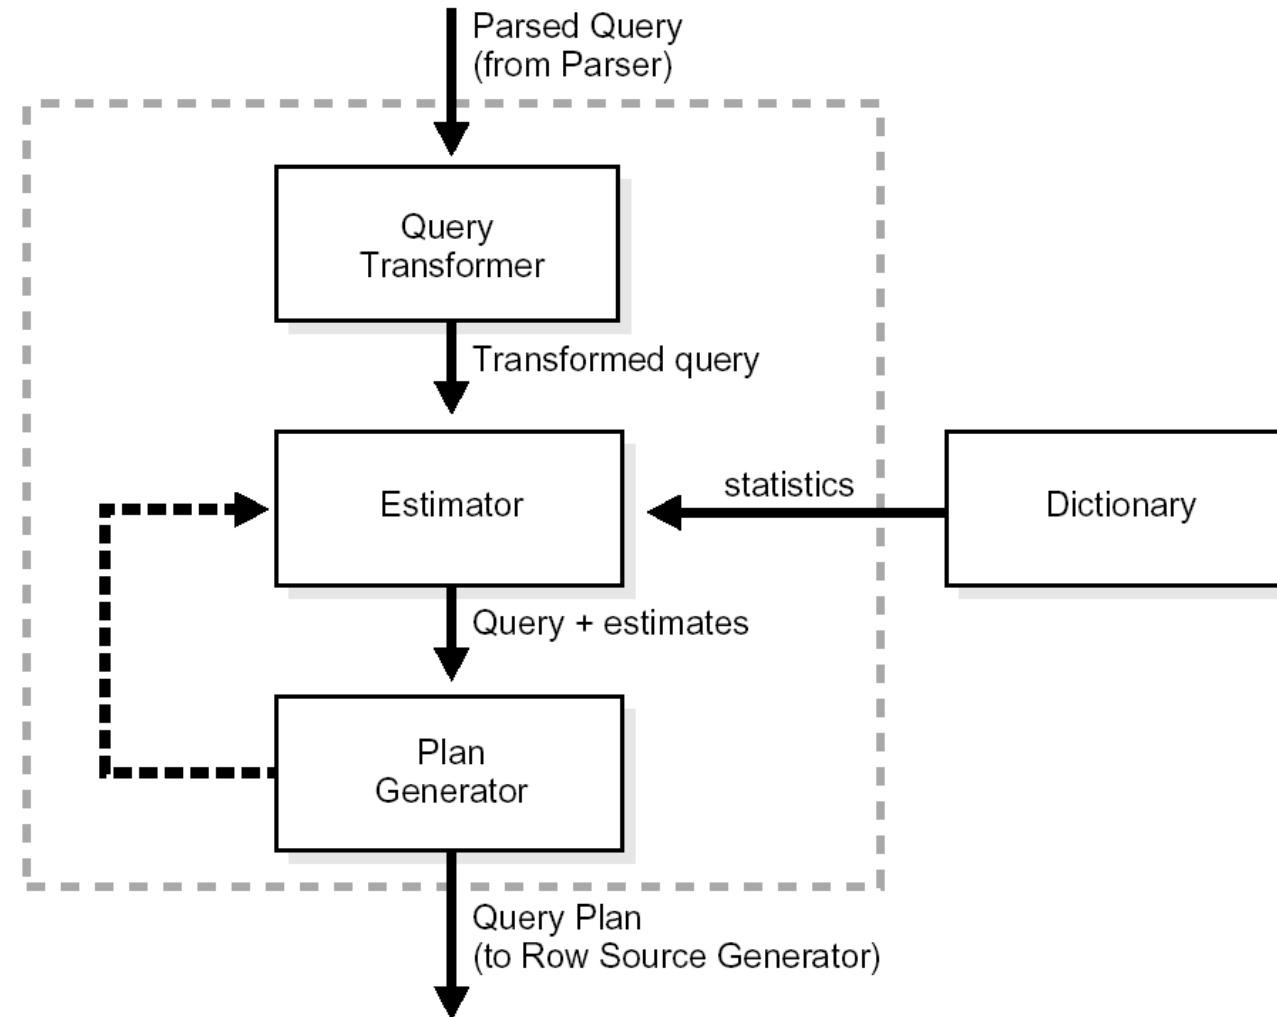
# Kardinalität und Selektivität

$$\textit{Cardinality} = \textit{Selectivity} \cdot \textit{TotalNumberOfRows}$$

## Cardinality is derived from

- Table statistics (total number of rows)
- Column statistics (number of distinct values, histograms)
- WHERE condition

# Architektur des Oracle Optimizers



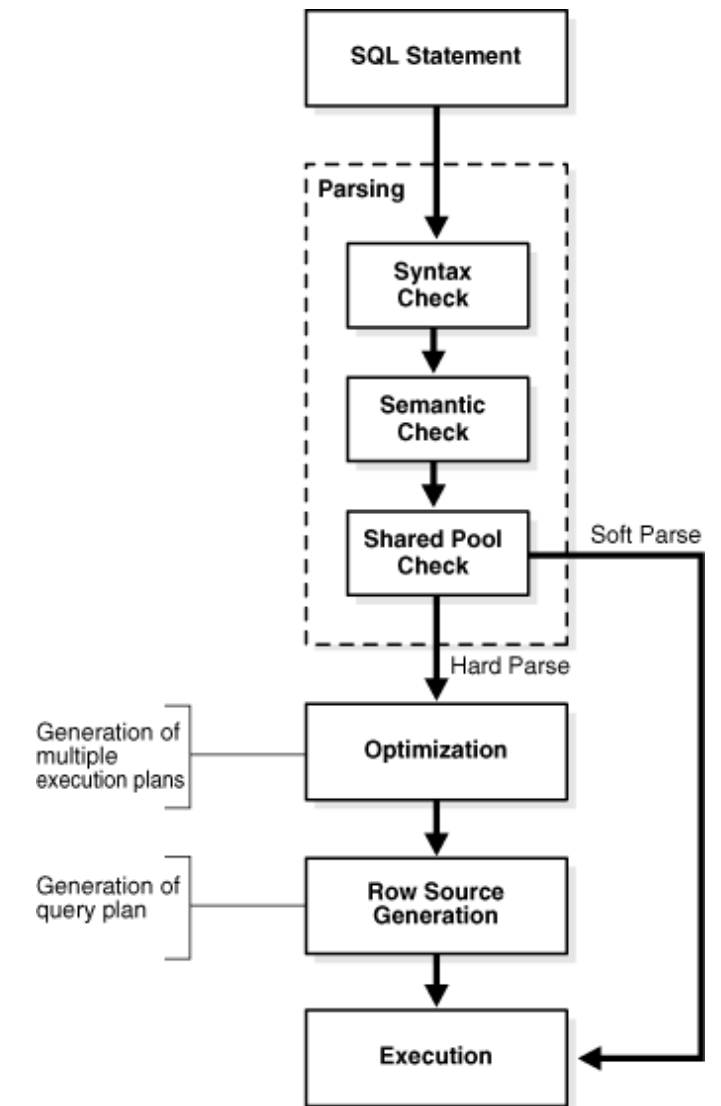
# Parsing

## Optimierung wird während Parse-Phase durchgeführt

- Parser überprüft SQL-Statement
- Optimizer berechnet günstigsten Execution Plan
- SQL und Plan werden in Shared SQL Area (Shared Pool) gespeichert

Bei erneuter Ausführung des SQL-Befehls wird Soft Parse ausgeführt

- SQL im Shared Pool vorhanden,  
=> gleicher Execution Plan wird verwendet






# JOIN-METHODEN

Performanceoptimierung


# Join Trees – Deep Trees




Id	Operation	Name
1	HASH JOIN	
2	HASH JOIN	
3	HASH JOIN	
4	TABLE ACCESS FULL	T1
5	TABLE ACCESS FULL	T2
6	TABLE ACCESS FULL	T3
7	TABLE ACCESS FULL	T4

Id	Operation	Name
1	HASH JOIN	
2	TABLE ACCESS FULL	T1
3	HASH JOIN	
4	TABLE ACCESS FULL	T2
5	HASH JOIN	
6	TABLE ACCESS FULL	T3
7	TABLE ACCESS FULL	T4

## Right Deep Tree




# Join Trees – Zig-Zag and Bushy Trees




Id	Operation	Name
1	HASH JOIN	
2	HASH JOIN	
3	TABLE ACCESS FULL	T1
4	HASH JOIN	
5	TABLE ACCESS FULL	T2
6	TABLE ACCESS FULL	T3
7	TABLE ACCESS FULL	T4

Id	Operation	Name
1	HASH JOIN	
2	VIEW	
3	HASH JOIN	
4	TABLE ACCESS FULL	T1
5	TABLE ACCESS FULL	T2
6	VIEW	
7	HASH JOIN	
8	TABLE ACCESS FULL	T3
9	TABLE ACCESS FULL	T4



# Nested Loops Join

- The left input (outer loop) is executed only once. The right input (inner loop) is potentially executed many times
- They are able to return the first row of the result set before completely processing all rows
- They can take advantage of indexes to apply both restrictions and join conditions
- They support all types of joins



# Nested Loops Join – Beispiel

```

SELECT /*+ ordered use_nl(t2 t3 t4) */ t1.*, t2.*, t3.*, t4.*
FROM t1, t2, t3, t4
WHERE t1.id = t2.t1_id
AND t2.id = t3.t2_id
AND t3.id = t4.t3_id AND t1.n = 19
  
```

Id	Operation	Name	Starts	E-Rows	A-Rows
0	SELECT STATEMENT		1		1000
1	NESTED LOOPS		1		1000
2	NESTED LOOPS		1	1000	1000
3	NESTED LOOPS		1	100	100
4	NESTED LOOPS		1	10	10
5	TABLE ACCESS BY INDEX ROWID	T1	1	1	1
* 6	INDEX RANGE SCAN	T1_N	1	1	1
7	TABLE ACCESS BY INDEX ROWID	T2	1	10	10
* 8	INDEX RANGE SCAN	T2_T1_ID	1	10	10
9	TABLE ACCESS BY INDEX ROWID	T3	10	10	100
* 10	INDEX RANGE SCAN	T3_T2_ID	10	10	100
* 11	INDEX RANGE SCAN	T4_T3_ID	100	10	1000
12	TABLE ACCESS BY INDEX ROWID	T4	1000	10	1000

# Block Prefetching

## Nested Loops Join without Block Prefetching

	Id	Operation	Name
	0	SELECT STATEMENT	
	1	NESTED LOOPS	
	2	TABLE ACCESS BY INDEX ROWID	T1
*	3	INDEX UNIQUE SCAN	T1_N
	4	TABLE ACCESS BY INDEX ROWID	T2
*	5	INDEX RANGE SCAN	T2_T1_ID

```
3 - access("T1"."N"=19)
```

```
5 - access ("T1"."ID"="T2"."T1 ID")
```

## Nested Loops Join with Block Prefetching

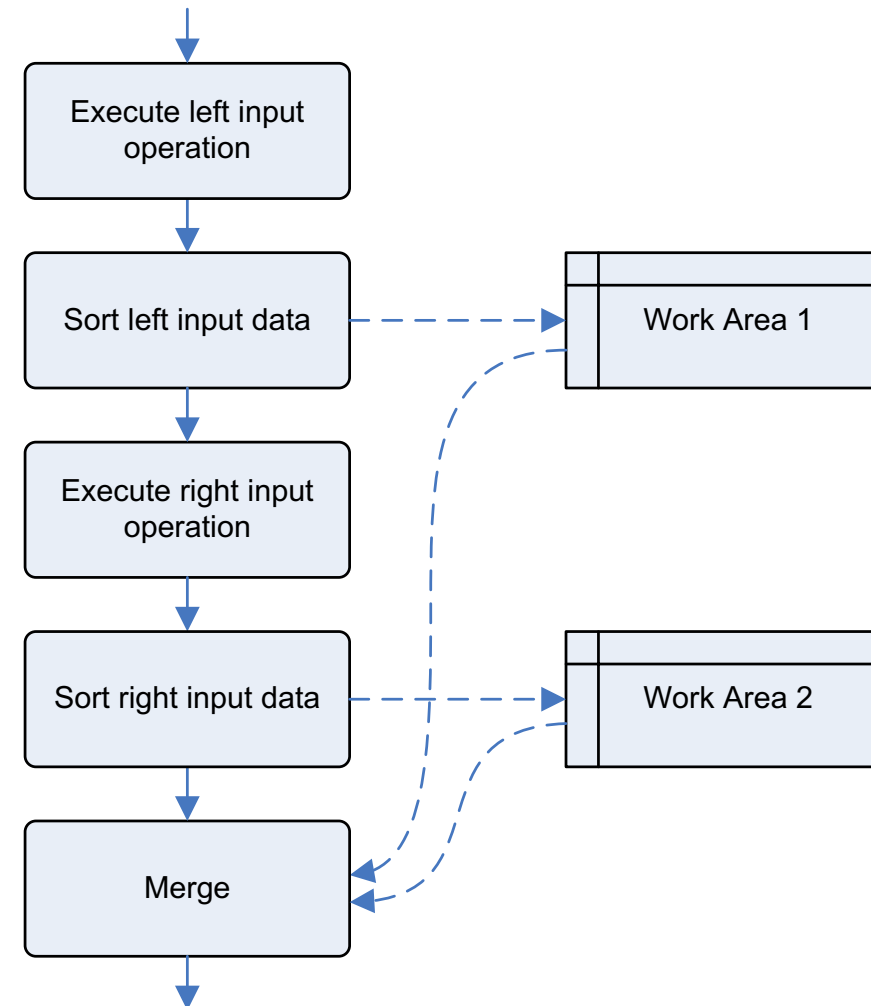
	Id	Operation	Name
	0	SELECT STATEMENT	
	1	NESTED LOOPS	
	2	NESTED LOOPS	
	3	TABLE ACCESS BY INDEX ROWID	T1
*	4	INDEX RANGE SCAN	T1_N
*	5	INDEX RANGE SCAN	T2_T1_ID
	6	TABLE ACCESS BY INDEX ROWID	T2

```
4 - access ("T1"."N"=19)
```

```
5 - access("T1"."ID"="T2"."T1 ID")
```

# Merge Join

- Each input is executed only once
- Both inputs must be sorted according to the columns of the join condition before returning the first row of the result set
- All types of joins are supported
- Sorting can be a very "expensive" operation → useful if data sources are already sorted



# Merge Join – Beispiel

```


SELECT /*+ leading(t1 t2 t3) use_merge(t2 t3 t4) */ t1.*, t2.*, t3.*, t4.*
FROM t1, t2, t3, t4
WHERE t1.id = t2.t1_id
AND t2.id = t3.t2_id
AND t3.id = t4.t3_id
AND t1.n = 19
  
```

Id	Operation	Name	Starts	E-Rows	A-Rows
0	SELECT STATEMENT		1		1000
1	MERGE JOIN		1	1000	1000
2	SORT JOIN		1	100	100
3	MERGE JOIN		1	100	100
4	SORT JOIN		1	10	10
5	MERGE JOIN		1	10	10
6	SORT JOIN		1	1	1
* 7	TABLE ACCESS FULL	T1	1	1	1
* 8	SORT JOIN		1	100	10
9	TABLE ACCESS FULL	T2	1	100	100
* 10	SORT JOIN		10	1000	100
11	TABLE ACCESS FULL	T3	1	1000	1000
* 12	SORT JOIN		100	10000	1000
13	TABLE ACCESS FULL	T4	1	10000	10000



# Hash Join

- Each input is executed only once
- The hash table is built on the left input only. Consequently, it is usually built on the smallest input
- Before returning the first row, only the left input must be fully processed
- Cross joins, theta joins, and partitioned outer joins are not supported



# Hash Join – Beispiel

## Left Deep Tree

```
SELECT /*+ leading(t1 t2 t3) use_hash(t2 t3 t4) */ t1.*, t2.*, t3.*, t4.*
FROM t1, t2, t3, t4
WHERE t1.id = t2.t1_id
AND t2.id = t3.t2_id
AND t3.id = t4.t3_id
AND t1.n = 19
```

Id	Operation	Name	Starts	E-Rows	A-Rows
0	SELECT STATEMENT		1		1000
* 1	HASH JOIN		1	1000	1000
* 2	HASH JOIN		1	100	100
* 3	HASH JOIN		1	10	10
* 4	TABLE ACCESS FULL	T1	1	1	1
5	TABLE ACCESS FULL	T2	1	100	100
6	TABLE ACCESS FULL	T3	1	1000	1000
7	TABLE ACCESS FULL	T4	1	10000	10000

# Hash Join – Beispiel

## Right Deep Tree

Id	Operation	Name	Starts	E-Rows	A-Rows
0	SELECT STATEMENT		1		1000
* 1	HASH JOIN		1	1000	1000
* 2	TABLE ACCESS FULL	T1	1	1	1
* 3	HASH JOIN		1	10000	10000
4	TABLE ACCESS FULL	T2	1	100	100
* 5	HASH JOIN		1	10000	10000
6	TABLE ACCESS FULL	T3	1	1000	1000
7	TABLE ACCESS FULL	T4	1	10000	10000

# Vergleich der Join-Methoden

